

Given the medical dataset that consist of age, sex, BMI(body mass index), children, smoker, and region as features, and charges as target variable, we have to apply linear regression to predict charges.

The hypothesis function looks like

$$h_{\theta}(x_i) = \theta_0 + \theta_1 \text{age} + \theta_2 \text{sex} + \theta_3 \text{bmi} + \theta_4 \text{children} + \theta_5 \text{smoker} + \theta_6 \text{region}$$

Tasks :

1. Load the dataset and do exploratory data analysis.

Exploratory data analysis is done to explore the data . The functions which I used to explore the data are `head()`, `info()`, `describe()`, `unique()` . To check whether there is any null value in the dataset, `df.isna().sum()` function is used .

2. Plot correlation between different variables and analyze whether there is a correlation between any pairs of variables or not.

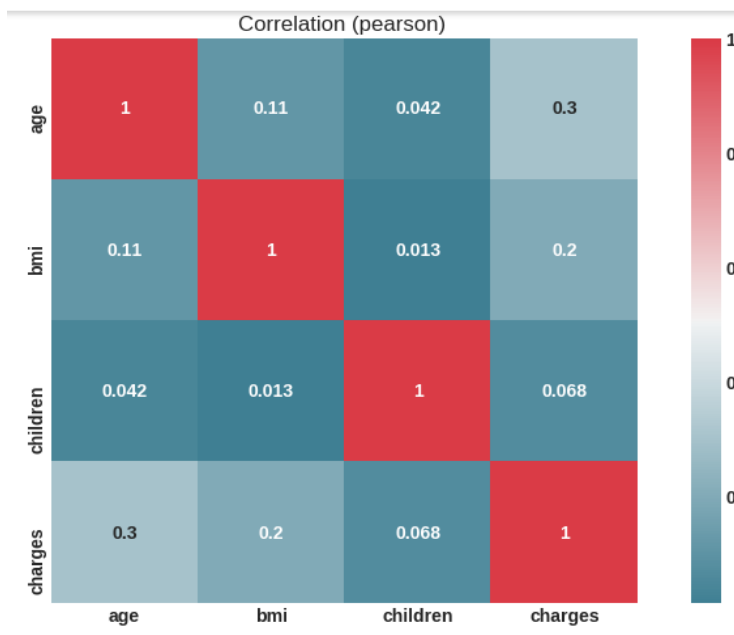
Pandas `df.corr()` is used to find the pairwise correlation of all columns in the dataframe. The function ignores non-numeric data type columns.

In the notebook, correlation matrix is created using two methods

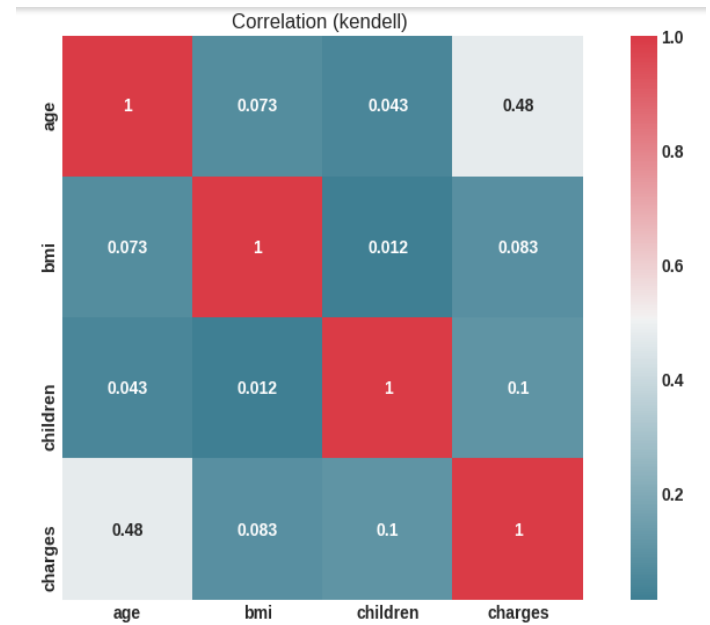
1. Pearson correlation
2. Kendall correlation

	age	bmi	children	charges
age	1.000000	0.109272	0.042469	0.299008
bmi	0.109272	1.000000	0.012759	0.198341
children	0.042469	0.012759	1.000000	0.067998
charges	0.299008	0.198341	0.067998	1.000000

	age	bmi	children	charges
age	1.000000	0.073273	0.043253	0.475302
bmi	0.073273	1.000000	0.011562	0.082524
children	0.043253	0.011562	1.000000	0.103107
charges	0.475302	0.082524	0.103107	1.000000


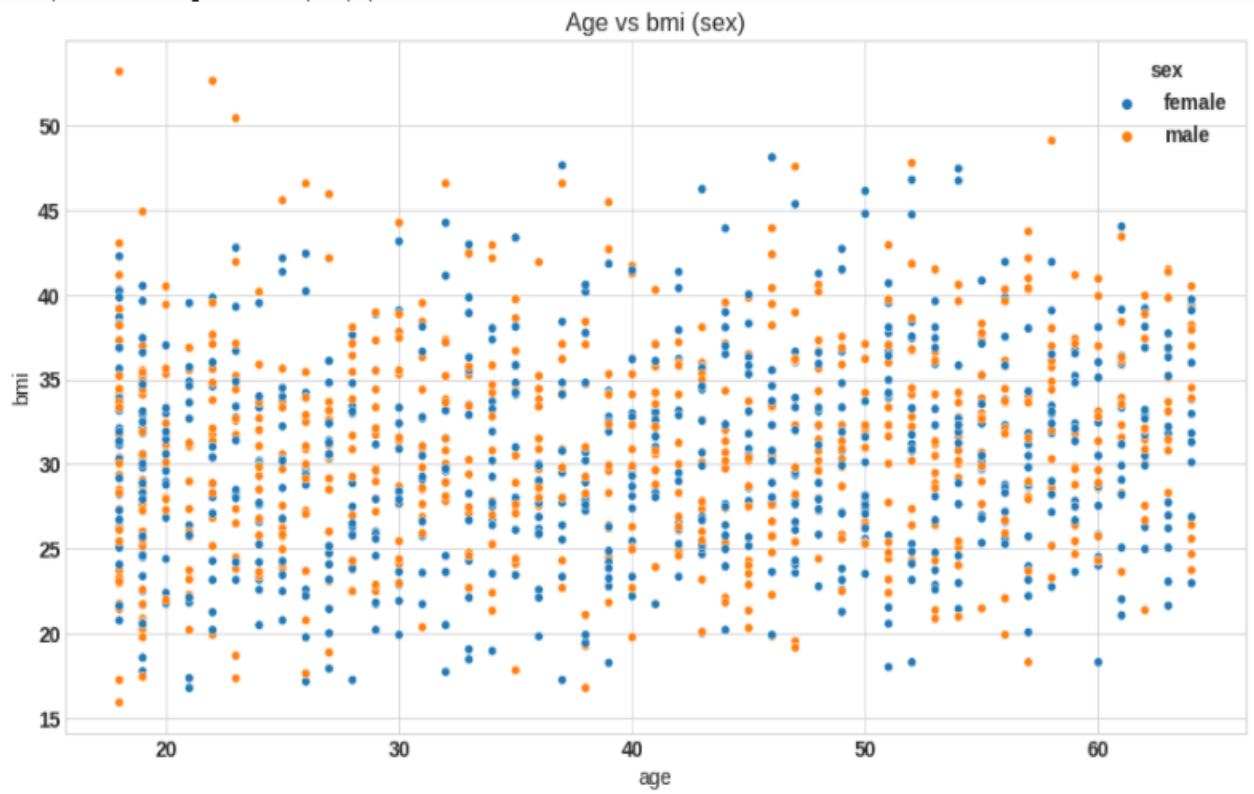


Pearson method




Kendall method

The graphs between pair of variables to get an idea of correlation:



3. Plot the distribution of the dependent variable and check for skewness (right or left skewed) in the distribution.

Before checking for skewness, we plot the boxplot model for 'Charges' column .



From the graph , it is apparent that there are many outliers in the data . Therefore to find those outliers we use z score method ,

```
z =  
np.abs(stats.zscore(df[['age', 'children', 'bmi', 'charges']]))
```

```
threshold = 3
```

```
print(np.where(z > threshold))
```

Removing the outliers

```
df= df[(z < threshold).all(axis=1)]
```

The number of samples remaining after removing the outliers are 1309

We have charges as the dependent variable . To check for skewness we can use ,


```
df['charges'].skew()
```

The output of this code is

```
Skewness:1.437160
```

It is clear from the value that the data is right skewed .

We can verify that using distplot:




4. Convert this distribution into normal by applying natural log and plot it. (If the distribution is normal then skip this).

The distribution can be converted into normal distribution using natural log of the values in charges column

```
df['log charges'] = np.log(df['charges'])
```

Plotting the distplot now :



We can do the same process using scipy boxcox function.

The resultant plot will be same for both methods

5. Convert categorical data into numbers. (You may choose one hot encoding or label encoding for that).

Label encoder is used to convert categorical data into numbers . The code can be seen in the notebook. The dataframe after categorial encoding looks like this:

	age	sex	bmi	children	smoker	region	charges	log charges	sp charges
0	19	0	27.900	0	1	3	16884.92400	9.734176	9.734176
1	18	1	33.770	1	0	2	1725.55230	7.453302	7.453302
2	28	1	33.000	3	0	2	4449.46200	8.400538	8.400538
3	33	1	22.705	0	0	1	21984.47061	9.998092	9.998092
4	32	1	28.880	0	0	1	3866.85520	8.260197	8.260197

Where log charges and sp charges are the column created while converting the distribution to normal

6. Split the data into training and testing sets with ratio 0.3.

```
X_train,X_test,Y_train,Y_test=train_test_split(x,y,test_size=0.3,random_state=42).
```

It is recommended to use the parameter `random_state=42` to produce the same results across a different run.

X_train has 916 samples and X_test has 393 samples

7. Build a model using linear regression equation $\theta=(X^T X)^{-1} X^T y$. (First add a feature $X_0=1$ to the original dataset).

To use the linear regression equation , we need to add a feature X_0 which will have values=1 for each data sample. We do this using :

```
X_train_s['X0'] = 1
```

```
X_test_s['X0'] = 1
```

Now to apply the equation, we will first convert the dataframe to numpy arrays .

```
X = X_train_s.to_numpy()
```

```
Y = Y_train.to_numpy()
```

The code below finds the theta using the linear regression equation

```
a = np.matmul(X_trans, X)
```

```
ainv = np.linalg.inv(a)
```

```
b = np.matmul(ainv , X_trans)
```

```
theta = np.matmul(b, Y)
```

8. Build a linear regression model using the sklearn library. (No need to add $X_0=1$, sklearn will take care of it.)

Using sklearn's `LinearRegression()`, we can easily build a regression model .

```
reg = LinearRegression()
```

```
reg.fit(X_train, Y_train)
```

9. Get the parameters of the models you built in step 7 and 8, compare them, and print comparisons in a tabular form.

The below table shows the weight of each feature calculated in step 7 and step 8

	Model (from scratch)	Using sklearn
age	0.03415366	0.03415366
sex	-0.07787483	-0.07787483
bmi	0.01175335	0.01175335
children	0.10252069	0.10252069
smoker	1.54601666	1.54601666
region	-0.04802618	-0.04802618
intercept	7.08517019	7.085170193937947

From the table , it is clear that the parameters are same for both .

10. Get predictions from both the models

For model built using linear regression equation:


```
y_pred = np.matmul(X_test_s_ar, theta)
```

Using sklearn:


```
y_pred_sk = reg.predict(X_test)
```



11. Perform evaluation using the MSE of both models



	Model(from scratch)	Using sklearn
MSE	0.17027210742136664	0.1702721074213662


12. Plot the actual and the predicted values to check the relationship between the dependent and independent variables.

For model 1(implemented from scratch):









Similarly the plots for model made from sklearn can be seen in the notebook.

The plots for both models are same

